


Input Image:



Question:

Is there a chair in
the image?

Answer:

No.

LVLM response:

Yes, there is a chair in
the image.

Token Probability Across 32 Decoder Layers

